

EECE5640
High Performance Computing
Homework 4

*** Submit your work on Canvas in a single zip file.**

1. (40) On the Discovery Cluster, develop an MPI program that computes pi using the dartboard method on 4 nodes. You can find a discussion of this method in this YouTube video:

https://www.youtube.com/watch?v=M34TO71SKGk&ab_channel=PhysicsGirl

Vary the number of darts and the number of nodes used. Report on the execution time of your program and the accuracy of your implementation. Try to find the best implementation between the number of darts each node is handling and the execution time. Note, you will need to use a different partition than reservation since our class partition only has 2 nodes.

2. (50) Develop a parallel histogramming program using C/C++ and MPI. A histogram is used to summarize the distribution of data values in a data set. The most common form of histogramming splits the data range into equal-sized bins called classes. For each class, the number of data values in the data set that fall into that class are totaled. Your input to this program will be integers in the range 1-1000 (use a random number generator that first generates the numbers first). Your input data set should contain N integers, where N can be varied.

Perform binning on each node using MPI, and then perform a reduction on node 0 to combine your partial results. Run this on 1, 2 and 4 nodes on Discovery. Your program should print out the values that fall into each class, in ascending class order.

- a. For your binning, use each node to produce a histogram for $(1/\# \text{ classes})^{\text{th}}$ of the input data values. Experiment with varying the number of classes and different values for N. Plot your performance results.
 - b. For your binning, use each node to only produce values for a single class of the input data values. Experiment with varying the number of classes and different values for N. Plot your performance results.
 - c. Compare the performance results for a and b. Attempt to explain any differences.
3. (10) Performance analysis of MPI applications has been an active area of research. There have been many performance tools developed. Please identify two of these frameworks and compare and contrast the capabilities of the toolsets you have selected. Make sure to cite all your resources. Please do not copy text out of user guides when you discuss the frameworks.

4. (Extra) Part of your weekly reading included a paper titled “MPI on Millions of Cores.” Given that this paper was published in 2010 (9 years ago), can you comment on what changes have occurred since 2010 that could positively and/or negatively impact our ability to fully exploit parallelism on millions of cores? Many of the papers today discuss exascale computing. Select a recent paper on exascale-computing and compare/contrast the barriers identified in the two papers that impact our ability to achieve these milestones.

This problem is worth 25 points of extra credit for the undergraduates in the class and 15 points of extra credit for the graduate students.

* Written answers to the questions should be included in your homework 4 write-up in pdf format. You should include your C/C++ programs and the README file in the zip file submitted.